

Lakehouse Data Platform: Technical Design Document

Document Control

Status	Draft for review
Version	1.0
Last updated	2026-10-03
Author	Gaurav Garg
Related documents	Architecture diagram: Figure 1 , editable in lakehouse_architecture.drawio · Decision records: 00-decision-register.md · Test plan: tests/README.md · Spec for AI-assisted development: docs/spec · Runbooks: 00-runbook-index.md · Deployment: deploy/README.md

1. Executive Summary

A single lakehouse for lending, insurance and recharge, on Apache Iceberg with an AWS Glue catalog (Figure 1). Money a customer sees is read live from its service, and an optional Flink lane serves signals that can't wait.

Two rules hold throughout:

1. **No gold table becomes readable until it is proven correct and complete.**
2. **Every answer served from the lake or the fast path states its version, how fresh it is, and whether it was verified.**

Section 10 goes deep on three problems: partner file versions, completeness and correctness at zero tolerance, and (less obvious) a consistent cut across one database's tables. Each has working code and tests. The one-time migration is designed in section 7.2.

2. Background and Problem Statement

Data comes from service databases, event streams, partner files, third-party APIs and hand-maintained Ops sheets. Analysts and finance report on it, applications read from it, data scientists build features from it, and auditors trace numbers back to it. Today it is scattered and inconsistent, and nobody can say whether a table is current or correct.

3. Goals and Non-Goals

Goals

- One platform for every business and source type, rebuildable from bronze.
- Every published table is provably **correct** (tied to the pause), **complete** (every declared input arrived) and **current** (its version and freshness stated).
- Every consumer gets an engine that suits how it reads, and the same question gets the same answer everywhere.
- Personal data is protected from the first byte stored, with enforceable erasure.

Non-goals

- A detailed plan for moving consumers off today's tables. It is sketched in 7.2.
- Automating corrections after publishing (risk 5).
- Cross-region disaster recovery. It is designed only.
- ML training and serving, choosing a BI tool, streaming joins, and chargeback.
- Blocking unsafe ad hoc queries at query time (section 7.3).

4. Requirements

4.1 Functional Requirements

ID	Requirement
FR-1	Ingest database changes, app events, partner files, ops sheets and API data
FR-2	Tokenise personal data before anything stores it
FR-3	Keep current-state tables equal to the source under replay, reordering and deletes
FR-4	Keep every version of every partner file, readable "as known at" any point in five years
FR-5	Publish a gold table only when it is complete and reconciled
FR-6	Serve dashboards and chat, ad hoc SQL, app reads, features and audit queries
FR-7	Trace any published number to its source position or file
FR-8	Load existing history once, with no gap and no overlap with the live stream
FR-9	(Extension) Serve labelled real-time signals

4.2 Non-Functional Requirements

Category	Requirement
Scale	10k events/s peak, 500M file rows/day, ~800M rows/day in total, 50M customers, five years kept (~1.5 trillion rows)
Freshness	Bronze ≤2 min, silver ≤7 min, gold ≤20 min, online store ≤5 min after gold, fast path in seconds
Correctness	Zero tolerance: gold ties to the paise, or it isn't published
Durability	Kafka keeps 7 days; bronze keeps everything as received for five years
Availability	A failed check never takes a table offline. Its last good version keeps serving, labelled
Compliance	RBI, IRDAI, DPDP 2023; all data and compute in India
Auditability	Any month-end figure can be reproduced exactly, with lineage

5. Assumptions and Constraints

Area	Assumption
Platform	AWS, ap-south-1 only, for data localisation
System of record	The sources remain the system of record; the lake is derived. Section 7.4 depends on this
Money	Signed BIGINT paise. Rates in integer basis points
Partner files	A full snapshot per (partner, feed, posting date), possibly in parts, unless the contract says changes only. Disputes are raised within 14 days

Area	Assumption
Control totals	Each source reports its totals for a business date once that date closes
Business day	Asia/Kolkata. Each partner's cut-off is declared in its contract
App traffic	Peak reads in the low thousands per second, which favours Aurora over DynamoDB

6. High-Level Design

6.1 Architecture Overview




Figure 1. The platform in five layers, rendered from `lakehouse_architecture.mmd` by `make diagram`.




Figure 2. The detailed architecture, drawn by hand in draw.io; the editable source is `lakehouse_architecture.drawio`.

Continuous sources arrive through Kafka (database changes through Debezium on Kafka Connect, app events through an event collector) and bounded ones through S3; gold passes the publish gate before serving. Two lanes bypass the lake: money read live from its owning service, and the Flink fast path. Lineage and check results go to OpenMetadata; metrics to VictoriaMetrics and Grafana.

6.2 Data Layers (Medallion)

Layer	Holds	To get in, a record must
Bronze	Everything as received, append-only, with personal data already tokenised	Nothing is refused: unparseable rows land flagged for reprocessing, never dropped
Silver	Current state from database changes, file facts with full history, deduplicated events	Match its contract, be applied in order, and be deduplicated
Gold	Business tables, aggregates and features	Pass the gate: complete for its declared inputs, and tied to the paise

6.3 Technology Choices and Alternatives Considered

The critical choices have their own decision records. All the others are in the `decision register`.

Area	Chose	What it buys	Rejected	What it costs
Table format	Apache Iceberg	Branches to build gold out of sight; time travel for audits; an open format	Delta Lake, Hudi	We run compaction and snapshot expiry
Catalog	AWS Glue, with published versions (01)	Managed, and native to AWS	Nessie, Polaris, S3 Tables	No multi-table commits; readers resolve a published version
Processing	Spark for batch; Flink only for the fast path (08)	One engine for merges, builds and backfill; seconds only where needed	One engine for everything	Two engines to run

Area	Chose	What it buys	Rejected	What it costs
Change ingestion	Debezium → Kafka → connector appends; Spark merges in order (03)	Changes that replay and stay in order per key	The connector writing to silver	Two hops of latency
Storage	S3, partitioned on time, bucketed on merge keys	~200 files a day per events table; merges touch only the buckets they need	Partitioning by business line × type × hour (>200,000 files a day per table)	Bucket counts are fixed per table layout
BI	A ClickHouse copy of published gold, with Cube (07)	Sub-second dashboards; one definition per metric	Querying Iceberg directly	A second copy, with no history
Applications	Money from the owning service; Aurora for the rest (06)	Money is always current; a new version swaps in with one transaction	DynamoDB first; reading money from the lake	Two read paths in every app

7. Detailed Design

7.1 Data Ingestion

Sources are grouped by how they behave, not by who owns them:

- **Continuous sources go through Kafka**, for replay, per-key ordering and decoupling.
- **Bounded sources skip Kafka**. They can already be re-read from S3.

Personal data is tokenised on the way in (section 8).

Source	Route	When it misbehaves
Postgres (each business)	Debezium → Kafka → connector appends to bronze → ordered merge into silver, up to the transaction cut (10)	Replays and reordering: ignored, since only newer changes apply. A transaction split across topics: waits until whole. Unchanged large column: keeps its value. Restore: forces a re-snapshot. Truncate: fails loudly. Stalled connector: slot lag is alerted and capped
App events	Event collector (schema check, tokenise) → Kafka → bronze → deduplicated in silver	Duplicates: removed on <code>event_id</code> . Unparseable: kept, flagged. Too late for a closed date: recorded as an exception
Partner files	S3 → File Loader: register, check, append, flip pointer (04)	Resend: treated as a duplicate. Correction: replaces the old version. Parts: published only when all have arrived. Truncated: held. Missing: alert when its window closes. Adjustments to earlier dates: posted today; closed periods are never rewritten
Ops sheets	Same as partner files	A change to the GL mapping needs human approval
Third-party APIs	Raw response saved to S3 first, then the File Loader	Retries: re-parse from S3. Records visible late: each pull re-reads a lookback window and deduplicates. Pages shifting: paged by key, never by offset. No completeness signal: a periodic full sweep
Reference files	Same checks, run on arrival, then a "reference updated" event	As partner files, plus the truncation guard

7.2 Initial Data Migration

Decision record: [09](#). The one-time load must meet the change stream at one point:

- **Create the replication slot first**, then export from a replica (Debezium's own snapshot would hold the slot open for hours).
- **Load the export tagged with the slot's starting position, then stream**. The merge's "strictly newer" rule absorbs any overlap.
- **Archived partner files replay oldest first**; files without sequence numbers load under a weaker, labelled legacy contract.
- **Past gold is rebuilt** and checked against the old finance reports.

Limit. Current-state tables hold only today's value. So "as known at" before migration day exists only where the source kept history tables.

Consumers move over gradually. Old and new run side by side, each metric's owner signs it off, and consumers cut over one at a time, dashboards first.

7.3 Data Quality, Freshness and Semantics

Correctness is enforced by the publish gate before any reader can see gold (section 10.2).

Completeness. Each gold table declares its inputs, and each input proves itself complete by its own signal (section 10.2). A late input holds only the tables that depend on it.

Freshness and versioning. Every published table carries a **version label**. It records:

- the positions it was built from, **one per source**, never collapsed into one number;
- when it was published;
- whether it was reconciled to the source, or only checked against silver.

Readers resolve published versions only.

Semantics.

- Each table has a **contract** in the Open Data Contract Standard format ([samples](#)): grain, keys, units and column meanings.
- **Cube defines each metric once**, for both the dashboards and the chat bot.
- **A contract check in CI** catches upstream drift that reconciliation can't see.
- **OpenMetadata** holds the glossary, the lineage and the gate's results.

Limit. Nothing blocks a wrong aggregation run directly on ClickHouse or Spark. The platform labels what it serves; it doesn't police queries.

7.4 Data Serving and Access

Consumer	Gets	Through	Freshness	Guarantee
Analysts and finance	Dashboards, chat bot, ad hoc SQL	Cube on ClickHouse; MCP bot via Cube; Spark	~20 min	Only versions that passed the gate. Dashboard and chat agree
Applications	Money	The owning service's API or read replica	Live	The source itself, never a copy
	Derived data (history, summaries, eligibility)	Read API on Aurora	~5 min after gold	One whole version per answer, with its label
	Fast signals	Fast-path Kafka topic	Seconds	Always verified: false

Consumer	Gets	Through	Freshness	Guarantee
Data scientists	Features	Spark on silver's history and gold	Daily	Point-in-time: nothing leaks from after the prediction date
Auditors	Any month-end figure, traced to source	Spark on a month-end tag; OpenMetadata	Five years	Tags kept five years. Lineage reaches the change-log position or file

One version of the truth. Copies load only whole, labelled versions, the Read API pins the versions behind each response, and fast-path signals are labelled unverified (01).

7.5 Real-Time Fast Path (Extension)

Fraud scoring can't wait 20 minutes (08). Flink reads only complete facts (app events, outbox events and reference updates, never raw change topics) and publishes signals labelled unverified, also stored in bronze. The cost: nothing verifies them, and live and reconciled numbers can disagree.

7.6 AI-Assisted Analytics

The chat bot (an MCP server) answers questions through Cube's governed metrics, never by writing its own SQL, so it gives the same number as the dashboards. The data catalog supplies the context: table descriptions, the business glossary, lineage and each table's contract.

- **Guardrails:** only published gold; personal data only as tokens; the asker's own access, enforced through Cube's security context and Lake Formation; a question no metric defines is declined and routed to an analyst.
- **Every answer** carries its version label and freshness.
- **An evaluation set** of known questions with known answers runs on every change to the Cube model.

8. Security, Privacy and Compliance

- **Tokenisation** (02): vault-issued random tokens, applied at every entry point before anything stores a value: the Debezium transform chain, the event collector for app events, and the File Loader. If the vault is down, ingestion stops rather than letting clear text through.
- **Encryption at rest.** SSE-KMS on every S3 prefix, with one key per layer (landing, bronze, silver, gold), so access can be revoked layer by layer. MSK, Aurora and ClickHouse volumes, and the vault's keys, are KMS-encrypted. Raw files sit in the landing zone for 14 days, the only clear values outside the sources and the vault.
- **Encryption in transit.** TLS 1.2 or later on every hop: MSK with TLS and IAM authentication, Debezium to Postgres with `sslmode=verify-full`, Spark to S3 over HTTPS, and Aurora and ClickHouse with TLS required.
- **Compression** happens before encryption: zstd for Iceberg's Parquet files and for Kafka producers, and ClickHouse's default LZ4 with ZSTD on large columns.
- **Access.**
 - Lake Formation grants per table.
 - Readers get silver and gold views, never staging.
 - Held branches are open to on-call only.
 - Real values come only from the vault, with a stated purpose, and are logged.

- **Erasure, by retention class** set in each field's contract. Without a legal hold, the customer's vault entry is deleted at once (its backups expire within the erasure window), so every token for them becomes meaningless. Under a retention obligation (loans, policies, KYC, watchlists), the entry is masked outside regulatory use and deleted when the period ends. Month-end figures stay reproducible, and every decision is logged.

9. Operational Considerations

9.1 Orchestration and Maintenance

Airflow schedules file loads, gold builds, their gates, and maintenance. The gate is a blocking task, never a step data flows through. Compaction and snapshot expiry run on **closed partitions only**, because compacting a live partition collides with the writer.

9.2 Monitoring and Alerting

Table lag behind sources; pause breaks (target zero, any break pages); open held branches; transaction-cut lag; replication-slot lag; files per partition; commit retries; and fast-path drift from gold.

9.3 Failure Modes and Recovery

Failure	Recovery
Kafka consumer down < 7 days	Resume from offsets; replay is safe
Down > 7 days, or a database restored	Re-snapshot with a new incarnation; sweep rows the snapshot didn't see
Gate fails	Held; the previous version keeps serving, labelled
Bad code deployed	Roll back the snapshot; rebuild from bronze
Vault down	Ingestion stops, by design; watch slot lag
Region lost	ap-south-2 recovery is designed, not built or rehearsed

On-call procedures for these failures are in the [runbooks](#).

9.4 Cost

Order of magnitude only, not modelled. Always-on compute (MSK, Kafka Connect, Spark, Flink) costs most, then ClickHouse and Aurora, then about 60 TB of storage. Month-end tags pin their files for five years, the one retention decision whose cost grows without a ceiling until then.

Tracking and alerts.

- Every resource carries cost-allocation tags: `business_line`, `layer` and `component`.
- The AWS Cost and Usage Report lands in S3 and is queried through Athena. A Grafana dashboard shows cost per business line, per layer and per pipeline, and cost per million rows published.
- AWS Budgets per account, and Cost Anomaly Detection monitors per service and per tag, alert through SNS into the same Alertmanager routes as operational alerts.
- Signals watched: Spark hours, S3 request counts (a small-files warning), MSK and ClickHouse.

9.5 Storage Lifecycle and Archival

S3 lifecycle rules **never expire Iceberg data files**: deleting files behind a table's back corrupts it. Lake retention runs through Iceberg itself (`expire_snapshots`, `remove_orphan_files`, and deleting partitions past retention). Lifecycle rules cover only the prefixes that aren't tables.

Data	Rule
Landing zone (raw files)	Expire after 14 days
Failed bucket	Expire after 90 days
Query results and job logs	Expire after 30 days
Kafka topics	7 days' retention
Iceberg tables (bronze, silver, gold)	S3 Intelligent-Tiering, with Archive Instant Access only, so a read never waits for a restore
Month-end tags	Kept five years; exported to Glacier Deep Archive only where regulation needs longer

10. Key Technical Challenges (Deep Dives)

How these were chosen. About 37 ways the sources misbehave were listed. A problem counts as hard only if it happens routinely, gives a plausible wrong answer that passes natural checks, has no off-the-shelf fix, and affects trusted numbers. Out-of-order changes to one key didn't qualify: the ordered merge (section 7.1) is standard practice. Nor did incremental API pulls: the discipline in section 7.1 covers them. Three larger problems are named in section 14.

10.1 Partner Files: Which Version Is the Truth, and Is It Whole?

Why it's hard. A delivery (partner, feed, posting date) can arrive in parts, twice, corrected, older-after-newer, or truncated. Each part can tie to its own trailer while the delivery is wrong.

Guarantee. A version publishes only when all its parts have arrived. An older version never wins, however publishes race. No part loads twice, even after a crash. A version far smaller than the same weekday usually is gets held. "As known at" works for any past moment.

Design (04, [code](#)). Register each part with its version and place in the manifest; check it whole; load it by replacing any rows with its file ID, so reruns and racing workers load it once; then flip the current-version pointer by compare-and-set, decided inside one MERGE.

Trade-offs. Readers go through the pointer, every version is stored, and a new feed has no baseline for the truncation guard.

10.2 Completeness and Correctness at Zero Tolerance

Why it's hard. Each source proves completeness differently, and keeps moving while we compare:

Source	Proof of completeness
Partner files	A current version of the delivery, and the arrival calendar
Databases	The transaction cut has passed the date's cut-off (section 10.3)
APIs	A periodic full sweep; otherwise labelled unverifiable
App events	None, so events never feed money

Guarantee. No gold version is readable unless every input is complete by its own signal and every (business line, date) ties exactly on count, signed paise, distinct keys and a content fingerprint.

Design (05, [code](#)). Build on a branch, replacing the whole date window. Compare with a full outer join against a **pinned** reference: silver at the snapshot the build read, not silver now, and the source's totals after close. Only the fingerprint catches ₹5,000 moved between customers. It is a 32-bit hash per row, summed, defined identically for Postgres and Spark, so it cannot overflow. Publish by `fast_forward`, or hold.

Trade-offs. Sources must report daily totals, and the intraday check can't see a change lost before silver.

10.3 A Consistent Cut Across One Database (Hidden)

Why it's hidden. One transaction (a loan and two ledger lines) lands as events on different partitions at different moments. Silver briefly holds the loan alone, and gold built then **ties to silver** because silver itself is torn. Idle and lagging partitions also look the same, so "is the day complete?" has no answer.

Guarantee. Gold never reads a half-applied transaction, and a date is complete only when every transaction before its cut-off is in silver.

Design (10, [code](#)). Debezium's transaction records give each transaction's event count per table. The cut is the longest prefix of transactions whose events have all landed. Every table of the source merges up to it and is tagged with it, and gold reads all inputs at one tag.

Trade-offs. Silver is as fresh as the slowest partition, and one failed merge holds the source.

11. Testing and Development

Every test provokes a plausible **wrong answer**, not a crash, and asserts the specific wrong result a broken design would produce.

- Guarantees that rest on engine behaviour run against a real Iceberg engine: 31 tests, a few minutes.
- **Mutation testing:** each of 20 guards is broken in turn, and a test must fail. All 20 are caught.

The [test plan](#) maps each test to the wrong answer it prevents, lists the cases planned but not yet written, and explains how to run them.

AI-assisted, spec-driven development. The repository is built to be developed with an AI coding assistant working from the [spec](#):

- a **constitution** of rules no change may break;
- **user stories** whose acceptance criteria are test IDs;
- a **technical plan** that records the engine behaviour to respect;
- a **task backlog**.

Each change starts as a plan the author reviews against the constitution. It is built test first, adds a mutation for any new guard, and updates the spec in the same change. CI (lint, tests, mutation testing, secret and dependency scans) gates every merge, and [CLAUDE.md](#) points the assistant to the spec.

12. Risks and Mitigations

Where it breaks first. Silver's merge at peak. Merge-on-read delete files pile up faster than compaction clears them, merges slow down, and the transaction cut, which waits for every table of a source, falls behind. Gold then holds and goes stale. It fails safe (stale and labelled, never wrong), and it is the first thing to load-test (risk 1).

#	Risk	Impact	Mitigation
1	Merge throughput and small files at peak (unmeasured)	Silver misses its freshness target	Load-test before go-live; compact closed partitions; if needed, move the merge to Flink with the same guards
2	One late input holds many tables	Dependent tables go stale	Per-table completeness; held tables keep serving, labelled
3	Metric definitions drift from schemas	Wrong numbers everywhere, invisible to reconciliation	The contract check in CI
4	A copy loses its version label	Stale numbers served with full confidence	Version-stamped copies; the Read API refuses unlabelled data; lag alerts
5	A correction after publishing	Must reach gold, every copy and the cache; the month-end tag can't change	A manual procedure (runbook): rebuild through the gate, reload the copies, and book the adjustment in the current month. Not automated (non-goal, section 3)

13. Open Questions

- **Schema registry:** Confluent works with today's configuration; Glue is native to AWS but needs that configuration changed.
- **Landing retention:** 14 days holds only if every partner's dispute window is 14 days or less.
- **App read traffic** is unmeasured. Above the low thousands per second, DynamoDB replaces Aurora ([06](#)).
- **Vault throughput at 10k events/s** is unmeasured. The fallback is in [02](#).

14. Honesty: Limitations, Uncertainties and AI Use

Deliberately left out. See Non-goals (section 3).

Not solved here: restatement of closed periods (a manual runbook); matching one bank credit to many transactions net of fees; and source purges that arrive as deletes.

Where I'm unsure: Debezium's transaction records and the Postgres fingerprint are taken from documentation and a reference implementation, not run; Cube's version-bound refresh on the deployed version; and Glue's commit behaviour versus the test catalog.

Status of the code. `src/` implements the three hard problems and the ordered merge. The plumbing around them is stubbed in `src/lakehouse/pipeline.py`, each stub stating its guarantee. Tests run on a real Iceberg engine, and mutation testing shows every guard is tested. The gate's checks are Spark functions a Great Expectations or Soda checkpoint would call; that wiring isn't built.

Use of AI.

I used Claude throughout: to draft documents, code and tests, to set out options and trade-offs, and to run probes against a real Iceberg engine. The decisions are mine, including several that went against its first suggestion:

- Iceberg with ClickHouse, and strict medallion layers;
- Cube with an MCP bot on gold;
- Glue over Nessie and Polaris;
- an append-only connector, after seeing what upserts lose;
- Great Expectations or Soda for the gate;
- Aurora over DynamoDB, and Spark over Trino for audits;
- a fast path for events and reference files;
- partner files as a hard problem, solved the standard way;
- choosing the hard problems from how the sources misbehave, not from the design;
- removing the semantic validator;
- the final architecture diagram, which I drew.

I checked the work by running it, not by reading it. Testing against a real engine, and deleting every guard to confirm a test fails, found bugs in code that had already been reviewed.

Appendix A: Glossary

Term	Meaning
Bronze, silver, gold	As received; merged and conformed; business-ready and verified
Change-log position	A database's own ordering of its changes (the Postgres LSN), used to apply changes in order
The gate	The check a gold table passes on a temporary branch before readers can see it
Held branch	A gold build that failed its check: kept 14 days for investigation, never published
Version label	On every published table and answer: source positions, when published, whether reconciled
Current-version pointer	Which file is current for each (partner, feed, date); it only flips to a newer sequence number
Transaction cut	The newest point up to which every transaction of one database has fully landed, across all its tables; silver advances only to it
Outbox event	A business fact written in the same database transaction as the change it describes
Tokenisation	Replacing personal data with a random, vault-issued token before anything stores it
Fast path	The Flink lane serving labelled, unverified signals in seconds

Appendix B: Sizing

Quantity	Working	Result
Events per day	10k/s peak ÷ 3 × 86,400	~290M
Rows per day	290M + 500M file rows	~800M
Rows kept	800M × 365 × 5	~1.46 trillion
Iceberg footprint	1.46 trillion × ~40 bytes (Parquet, zstd)	~58 TB
Kafka retention	~3 MB/s average × 7 days × 3 replicas	~5 TB
Kafka partitions	Sized for consumer parallelism; throughput alone needs 2	~540
Files per day (one events table)	24 hourly partitions × ~8 after compaction	~200